

CSEN 604: Databases II

Lecture 3

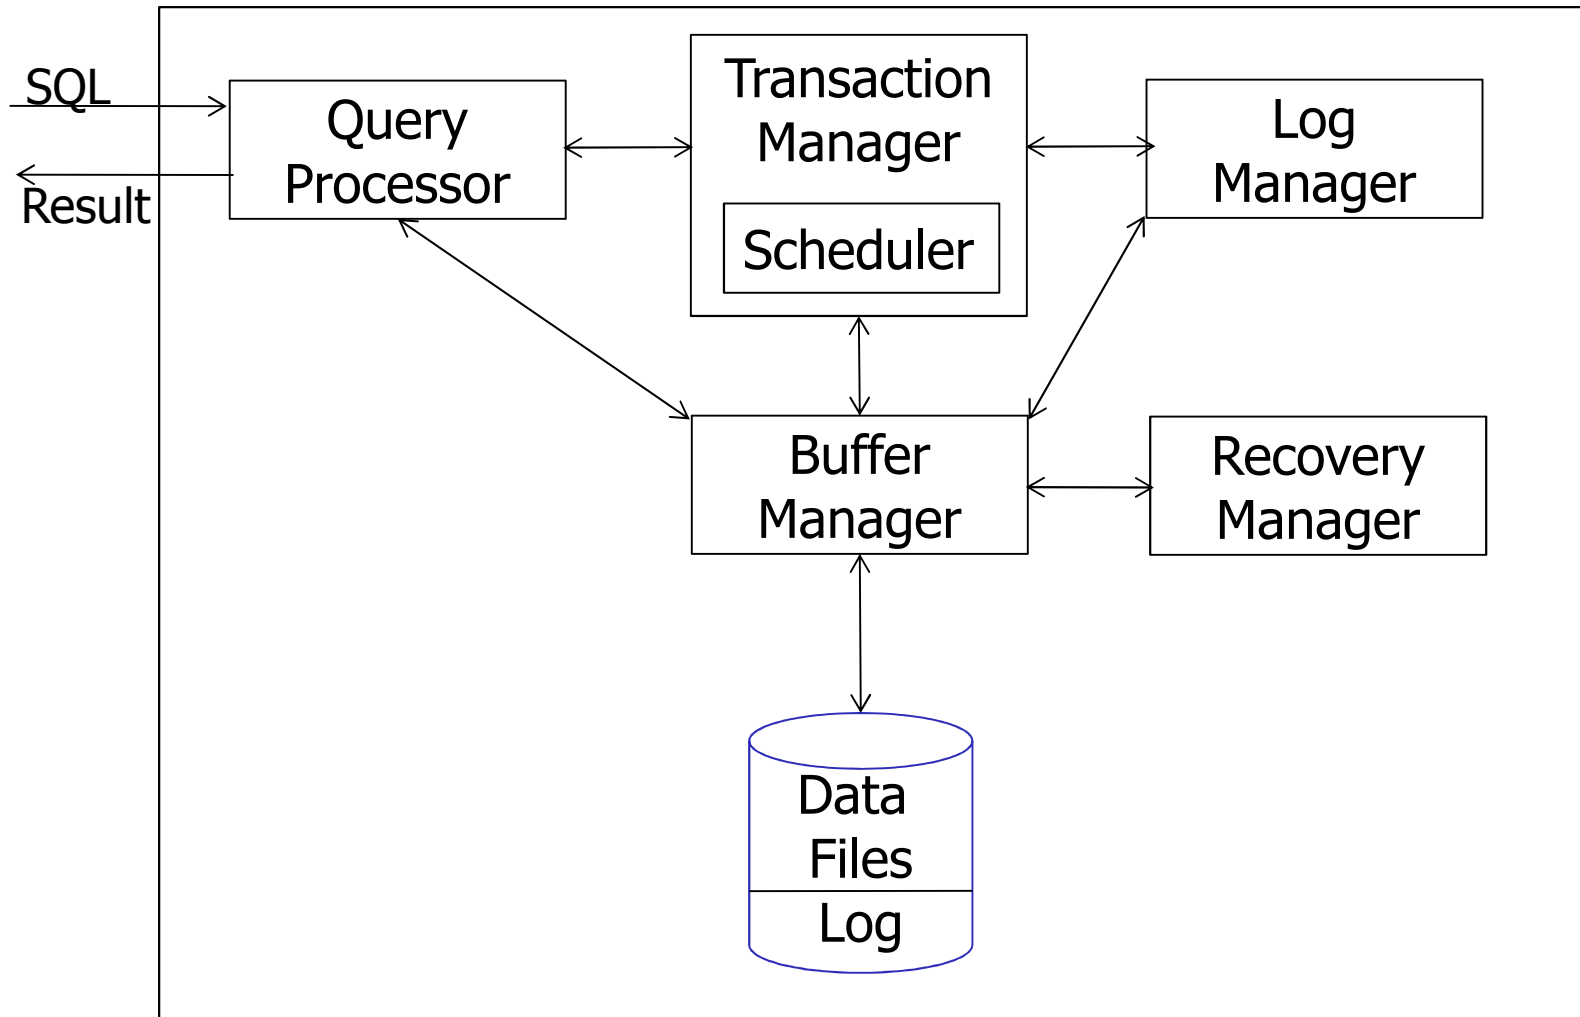
Dr. Wael Abouelsaadat
wael.abouelsaadat@guc.edu.eg
Office: C7.208

Office Hour is 4th slot Saturday or you can email for appointment

Acknowledgment: these slides are based on Prof. Garcia-Molina & Prof. Ullman slides accompanying the book: *Database Systems; the Complete Book*



DBMS Architecture





Topics

- B+ Tree
- Hashing



- B+ Tree
 - A data structure used to build an index
 - Most nodes are on hard-disk → a storage oriented data structures
 - Just like any tree: basic building blocks;
 - Node (link in CS3)
 - Node consists of data + references to children nodes.
 - Self-balanced data structure → $O(\log n)$



Deletion from B+tree

- (a) Simple case - no example
- (b) Coalesce with neighbor (sibling)
- (c) Re-distribute keys
- (d) Cases (b) or (c) at non-leaf




(b) Coalesce with sibling

- Delete 50

n=4

if $n = 4$, then **key** count
Non-leaf: max: 4, min: 2
Leaf: max: 4, min: 2






(b) Coalesce with sibling

- Delete 50

$n=4$

if $n = 4$, then **key** count
Non-leaf: max: 4, min: 2
Leaf: max: 4, min: 2






(c) Redistribute keys

- Delete 50

$n=4$

if $n = 4$, then **key** count
Non-leaf: max: 4, min: 2
Leaf: max: 4, min: 2






(c) Redistribute keys

- Delete 50

$n=4$

if $n = 4$, then **key** count
Non-leaf: max: 4, min: 2
Leaf: max: 4, min: 2






(d) Non-leaf coalesce

- Delete 37

$n=4$

if $n = 4$, then **key** count
Non-leaf: max: 4, min: 2
Leaf: max: 4, min: 2






(d) Non-leaf coalesce

– Delete 37

$n=4$

if $n = 4$, then **key** count
 Non-leaf: max: 4, min: 2
 Leaf: max: 4, min: 2




$n=4$

if $n = 4$, then **key** count

Non-leaf: max: 4, min: 2

Leaf: max: 4, min: 2



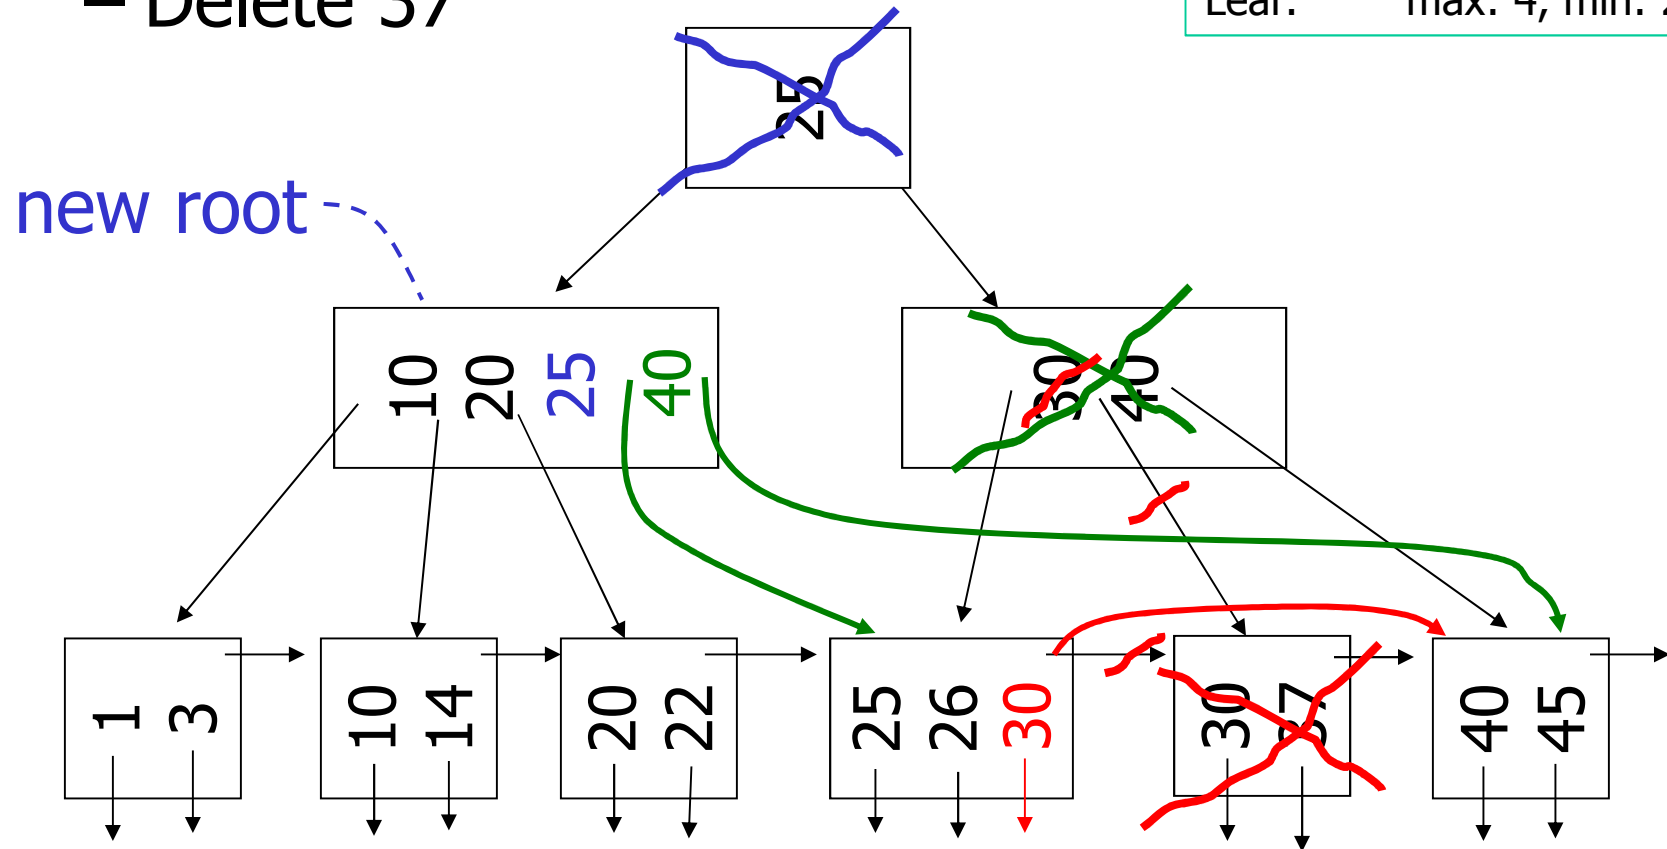


(d) Non-leaf coalesce

– Delete 37

$n=4$

if $n = 4$, then **key** count
 Non-leaf: max: 4, min: 2
 Leaf: max: 4, min: 2





When to create a B+ tree index ?

- You want to support efficient search on a single column

+

- A majority of received queries involve a range search (min, max, between, ...)



Creating an Index using B+ Tree

- PostgreSQL
 - SQL: `create index index_name on table-name(column-name) using B_tree;`
 - <http://www.postgresql.org/docs/8.2/static/sql-createindex.html>



Index vs. key vs. Clustering

- Key: logical concept
- Index: a secondary storage data structure
- Clustering: disk/page organizational issue
 - According to which column the data is stored in pages



Clustered Index vs. Non-clustered Index

- Clustered Index
 - Clustering defines the page content in terms of order tuples. Index and relation match in ordering
 - Only **one clustered index** can be created on a given database table.
 - Clustered indices can greatly increase overall speed of retrieval, but usually only where the data is accessed sequentially in the same or reverse order of the clustered index, or when a range of items is selected.



Clustered Index vs. Non-clustered Index

– Non-clustered Index


- logical ordering is specified by the index only.
- The physical order of the rows is not the same as the index order.
- The indexed columns are typically non-primary key columns (used in JOIN, WHERE, and ORDER BY clauses)
- There can be more than one non-clustered index on a database table.



Clustered Index vs. Non-clustered Index

– E.g. SQL Server


```
create table DemonstrationTable (  
  TableIdColumn int not null primary key nonclustered );
```






Hashing

$\text{key} \rightarrow h(\text{key})$



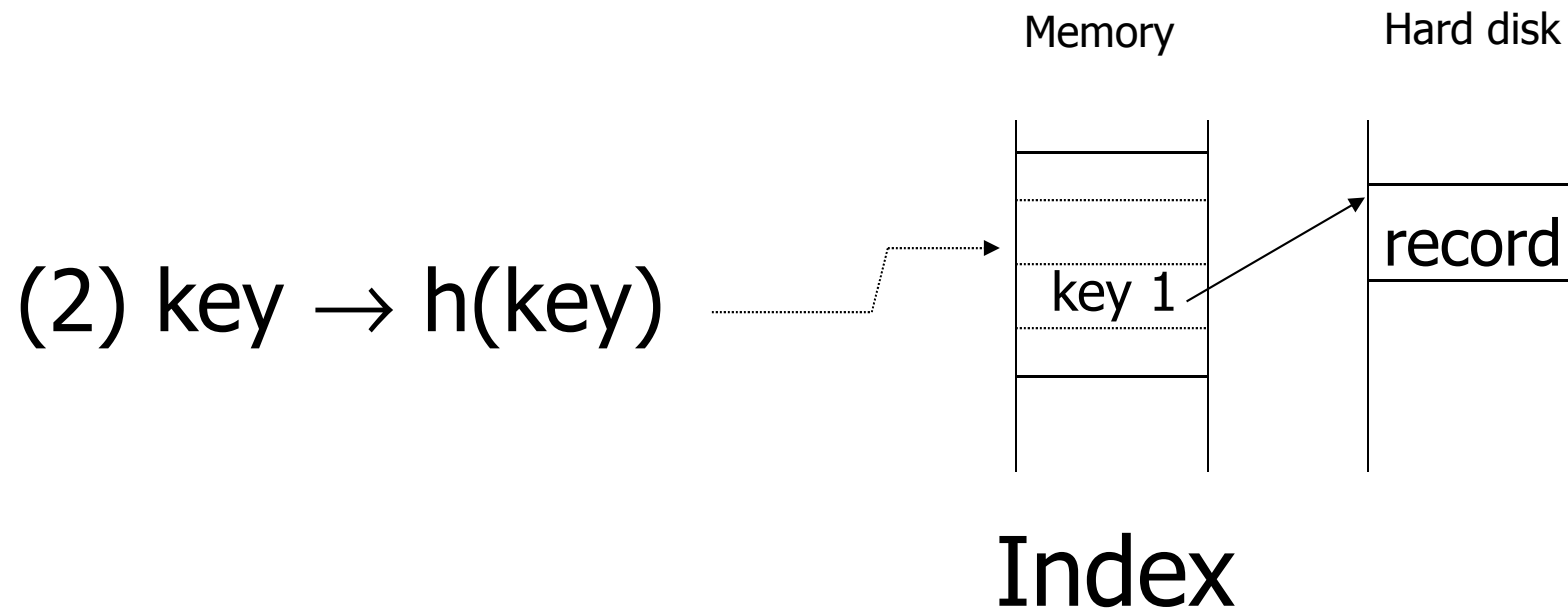


Two alternatives



- Alt (1) for “primary” search key

Two alternatives



- Alt (2) for “secondary” search key



Example hash function

- Key = ' $x_1 x_2 \dots x_n$ ' n byte character string
- Have b buckets
- h : add $x_1 + x_2 + \dots + x_n$
 - compute sum modulo b



- ➡ This may not be best function ...
- ➡ Read Knuth Vol. 3 if you really need to select a good function.



- ➡ This may not be best function ...
- ➡ Read Knuth Vol. 3 if you really need to select a good function.

Good hash
function:

➡ Expected number of
hash-value/bucket is the
same for all buckets




Within a bucket:

- Do we keep keys sorted?
- Yes, if CPU time critical
& Inserts/Deletes not too frequent



Next: example to illustrate inserts, overflows, deletes





EXAMPLE 2 records/bucket


INSERT:

$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$



→ This is the index, not the relation...



EXAMPLE 2 records/bucket

INSERT:


$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$

$h(e) = 1$





EXAMPLE 2 records/bucket

INSERT:

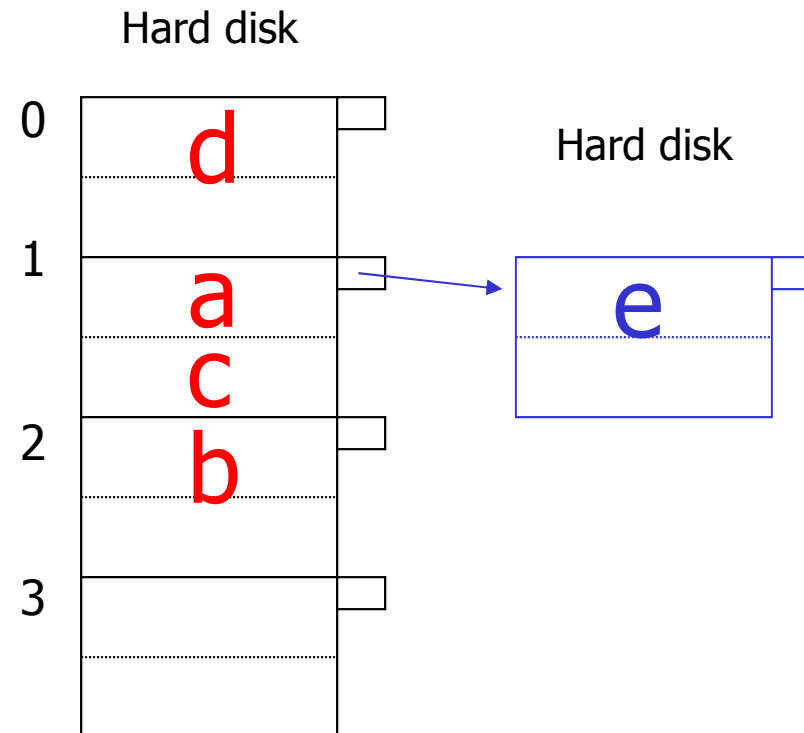
$h(a) = 1$

$h(b) = 2$

$h(c) = 1$

$h(d) = 0$

$h(e) = 1$

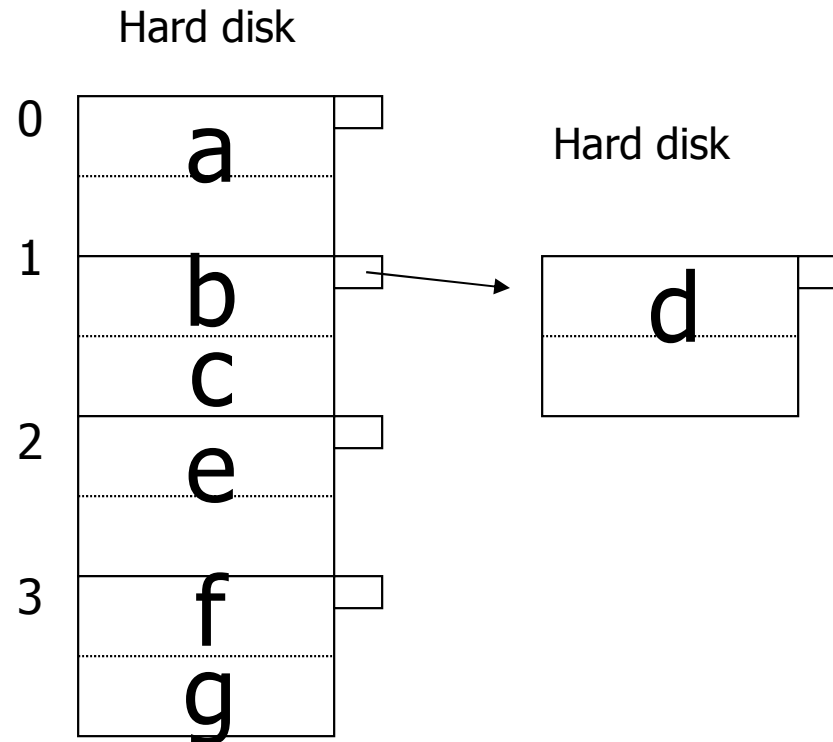




EXAMPLE: deletion

Delete:

e
f






EXAMPLE: deletion

Delete:

e

f

c





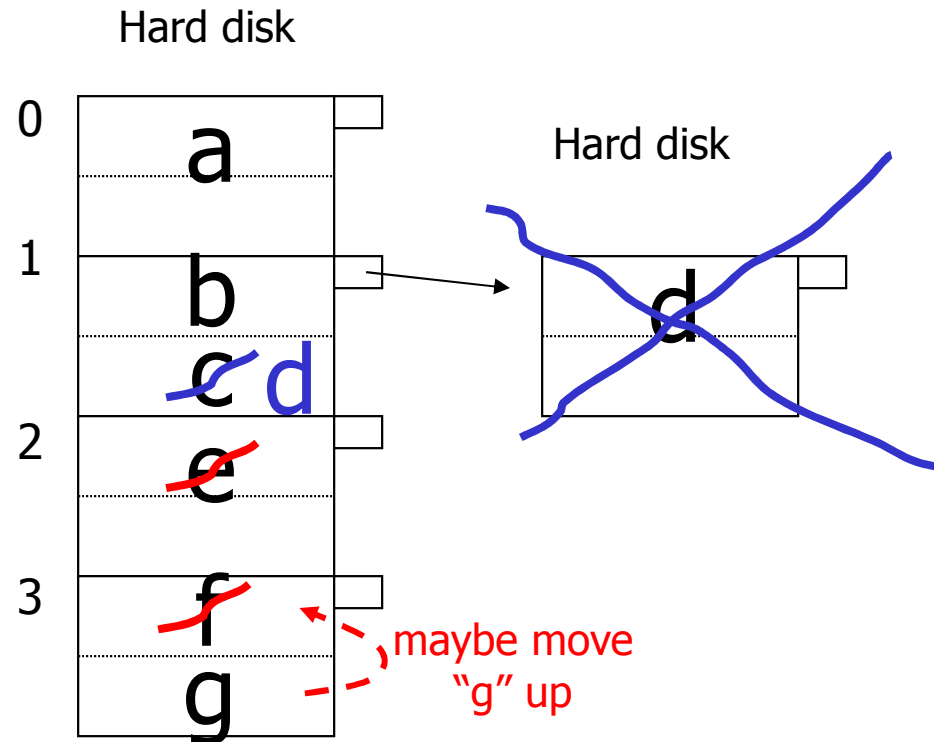
EXAMPLE: deletion

Delete:

e

f

c





Rule of thumb:

- Try to keep space utilization between 50% and 80%

$$\text{Utilization} = \frac{\text{\# keys used}}{\text{total \# keys that fit}}$$



Rule of thumb:

- Try to keep space utilization between 50% and 80%

$$\text{Utilization} = \frac{\text{\# keys used}}{\text{total \# keys that fit}}$$

- If $< 50\%$, wasting space
- If $> 80\%$, overflows significant
 ↖ depends on how good hash function is & on # keys/bucket



How do we cope with growth?

- Reorganizations
 - Very expensive → read all pages and rehash




How do we cope with growth?

- Reorganizations
 - Very expensive → read all pages and rehash
- Dynamic hashing
 - Extensible hashing
 - Linear hashing



Extensible hashing: two ideas


(a) Use i of b bits output by hash function




use $i \rightarrow$ grows over time....



(b) Use directory




Example: $h(k)$ is 4 bits; 2 keys/bucket




Insert 1010



Example: $h(k)$ is 4 bits; 2 keys/bucket





Insert 1010





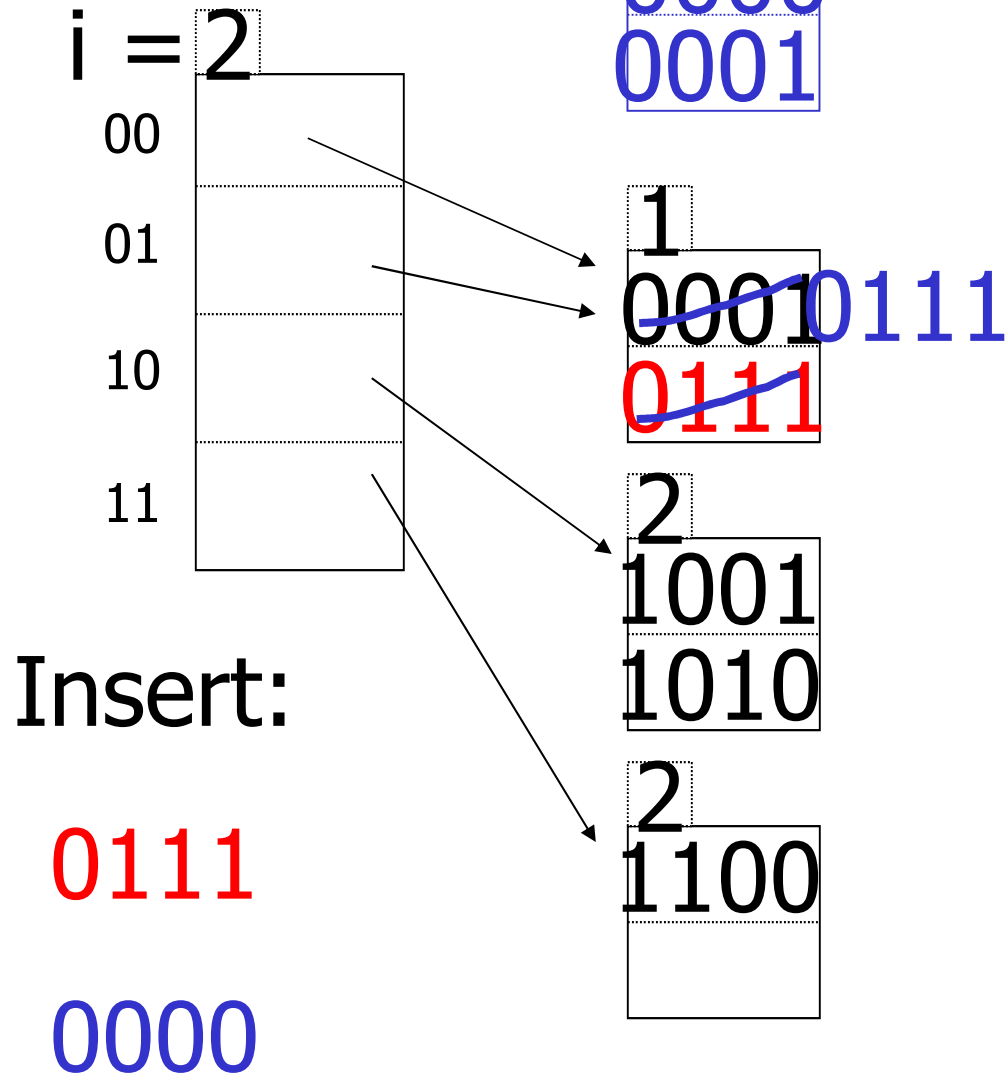
Example: $h(k)$ is 4 bits; 2 keys/bucket



Example continued

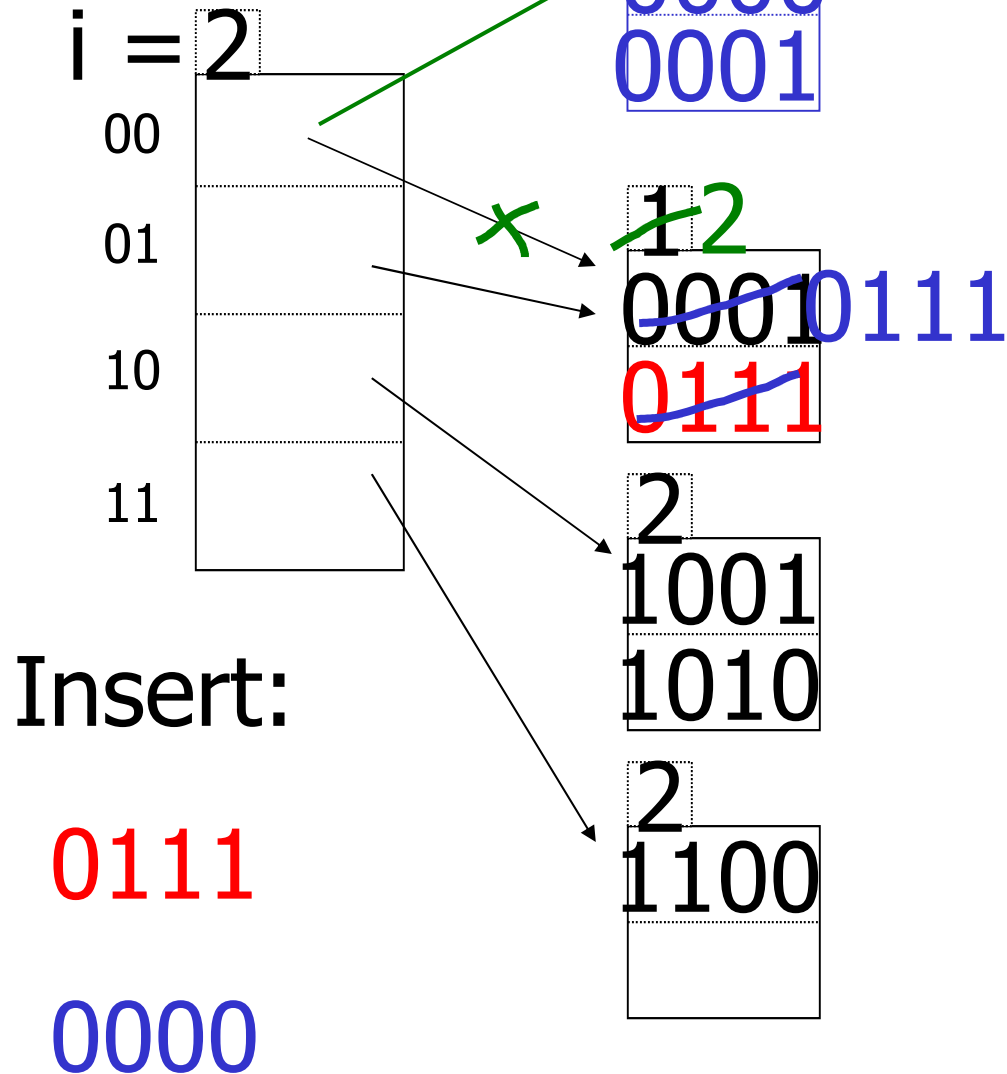


Example continued





Example continued





Example continued

$i = 2$

00

01

10

11

0000²
0001

0111²

1001²
1010


1100²

Insert:

1001

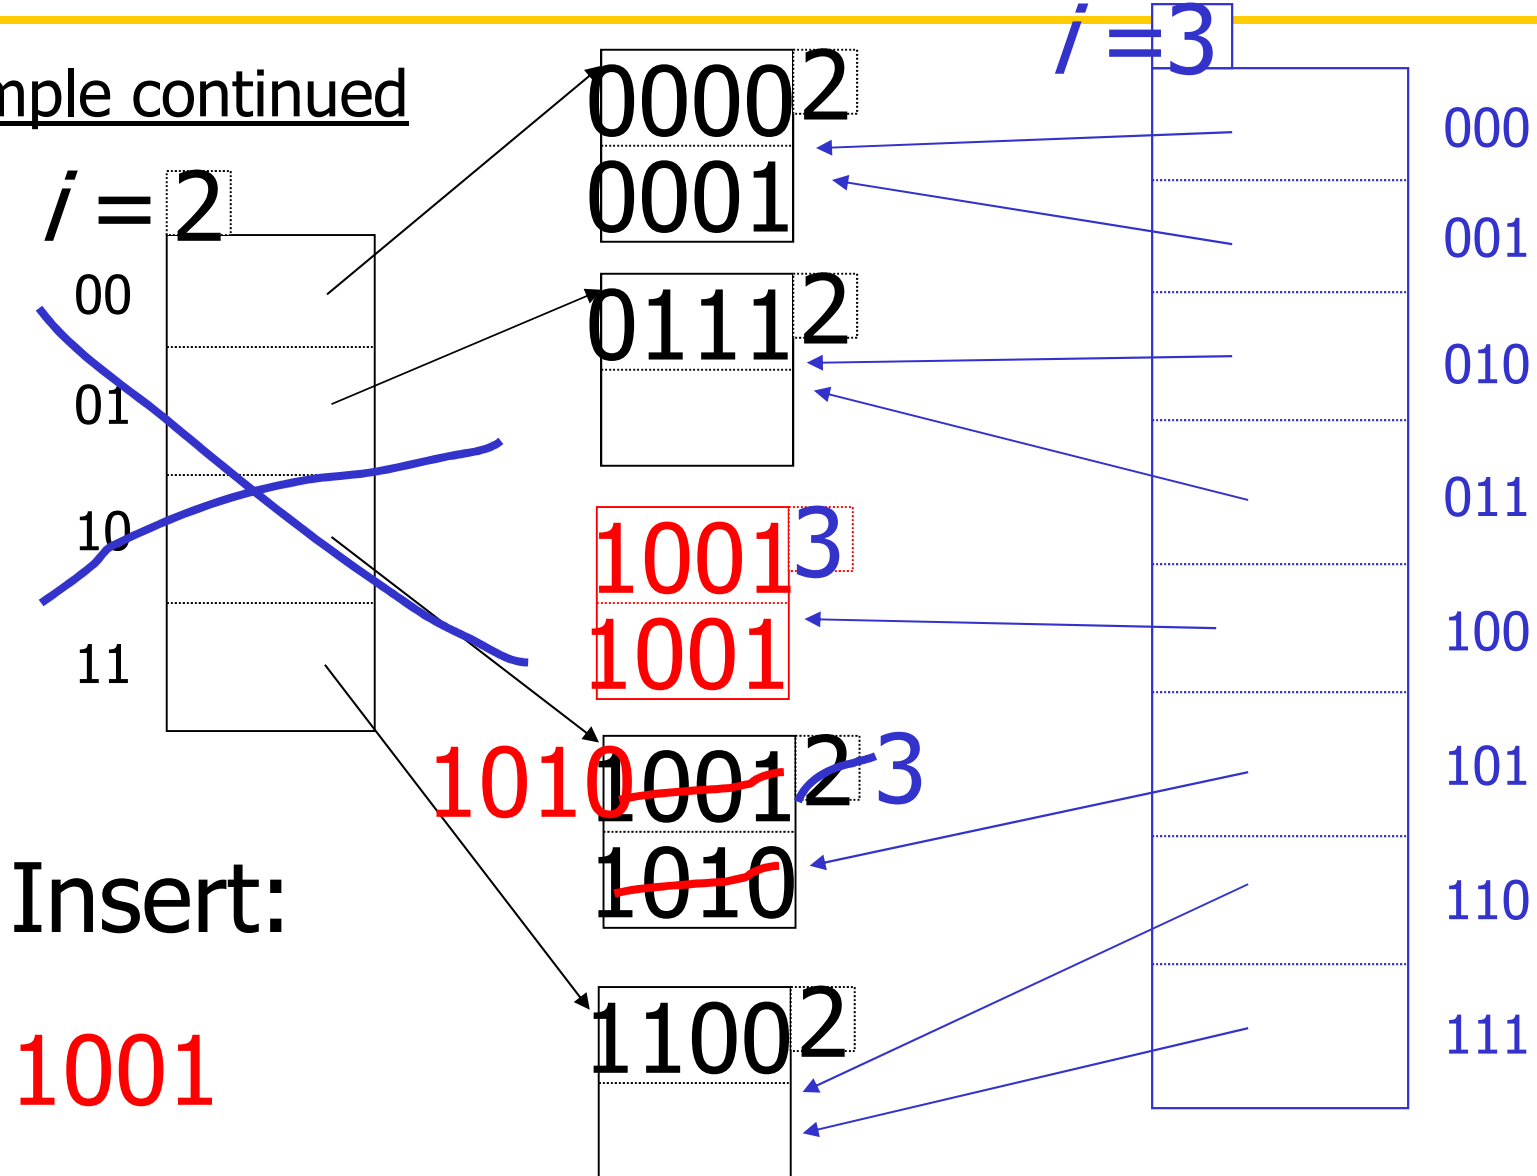


Example continued





Example continued





Extensible hashing: deletion

- No merging of blocks
- Merge blocks
and cut directory if possible
(Reverse insert procedure)



Deletion example:


- Run thru insert example in reverse!





Note: Still need overflow chains

- Example: many records with duplicate keys

insert 1100




if we split:






Solution: overflow chains

insert 1101



add overflow block:





Summary Extensible hashing

- ⊕ Can handle growing files
 - with less wasted space
 - with no full reorganizations



Summary Extensible hashing

- ⊕ Can handle growing files
 - with less wasted space
 - with no full reorganizations
- Indirection
 - (Not bad if directory in memory)
- Directory doubles in size
 - (Now it fits, now it does not)




Linear hashing

- Another dynamic hashing scheme

Two ideas:

(a) Use i low order bits of hash






Linear hashing


- Another dynamic hashing scheme

Two ideas:

(a) Use i low order bits of hash



(b) File grows linearly






☞ When do we expand file?

- Keep track of:
$$\frac{\text{\# used slots}}{\text{total \# of slots}} = U$$
- After every insertion, check if $U >$ threshold then increase buckets by 1
- If you run out of bits, add 1 more
 - 00 becomes 000




Example $b=4$ bits, $i=2$, 2 keys/bucket





Example $b=4$ bits, $i=2$, 2 keys/bucket



$m = 01$ (max used block)


Rule If $h(k)[i] \leq m$, then

look at bucket $h(k)[i]$
 else, look at bucket $h(k)[i] - 2^{i-1}$



Example $b=4$ bits, $i=2$, 2 keys/bucket

- insert 0101




$m = 01$ (max used block)

Rule If $h(k)[i] \leq m$, then


look at bucket $h(k)[i]$
 else, look at bucket $h(k)[i] - 2^{i-1}$



Example $b=4$ bits, $i=2$, 2 keys/bucket



- insert 0101
- can have overflow chains!




$m = 01$ (max used block)

Rule If $h(k)[i] \leq m$, then

look at bucket $h(k)[i]$
 else, look at bucket $h(k)[i] - 2^{i-1}$



Example $b=4$ bits, $i=2$, 2 keys/bucket




$m = 01$ (max used block)

- After every insertion, check if $U >$ threshold (e.g. 0.8) then increase buckets by 1

$$U = \frac{\text{\# used slots}}{\text{total \# of slots}}$$



Example $b=4$ bits, $i=2$, 2 keys/bucket






Example $b=4$ bits, $i=2$, 2 keys/bucket

0101

• insert 0101




$m = 01$ (max used block)
10

0101

-
- The diagram shows four blocks labeled 00, 01, 10, and 11. Each block contains a 2x2 grid of bits. Block 00 contains 0000 and 1010, with a red arrow pointing to the bottom row. Block 01 contains 0101 and 1111. Block 10 contains 1010 and is highlighted with a red border. Block 11 is empty and highlighted with a green border. A label 'Future growth buckets' with an arrow points to the green border of block 11. Below the blocks, the text $m = 01$ (max used block) is shown, with a red arrow pointing to the '01' and a green arrow pointing to the '11' block.



Example $b=4$ bits, $i=2$, 2 keys/bucket





Example Continued: How to grow beyond this?

$$i = 2$$

0000	0101	1010	1111
	0101		
00	01	10	11


...

$m = 11$ (max used block)



Example Continued: How to grow beyond this?

$$i = \cancel{2}3$$




$m = 11$ (max used block)



Example Continued: How to grow beyond this?

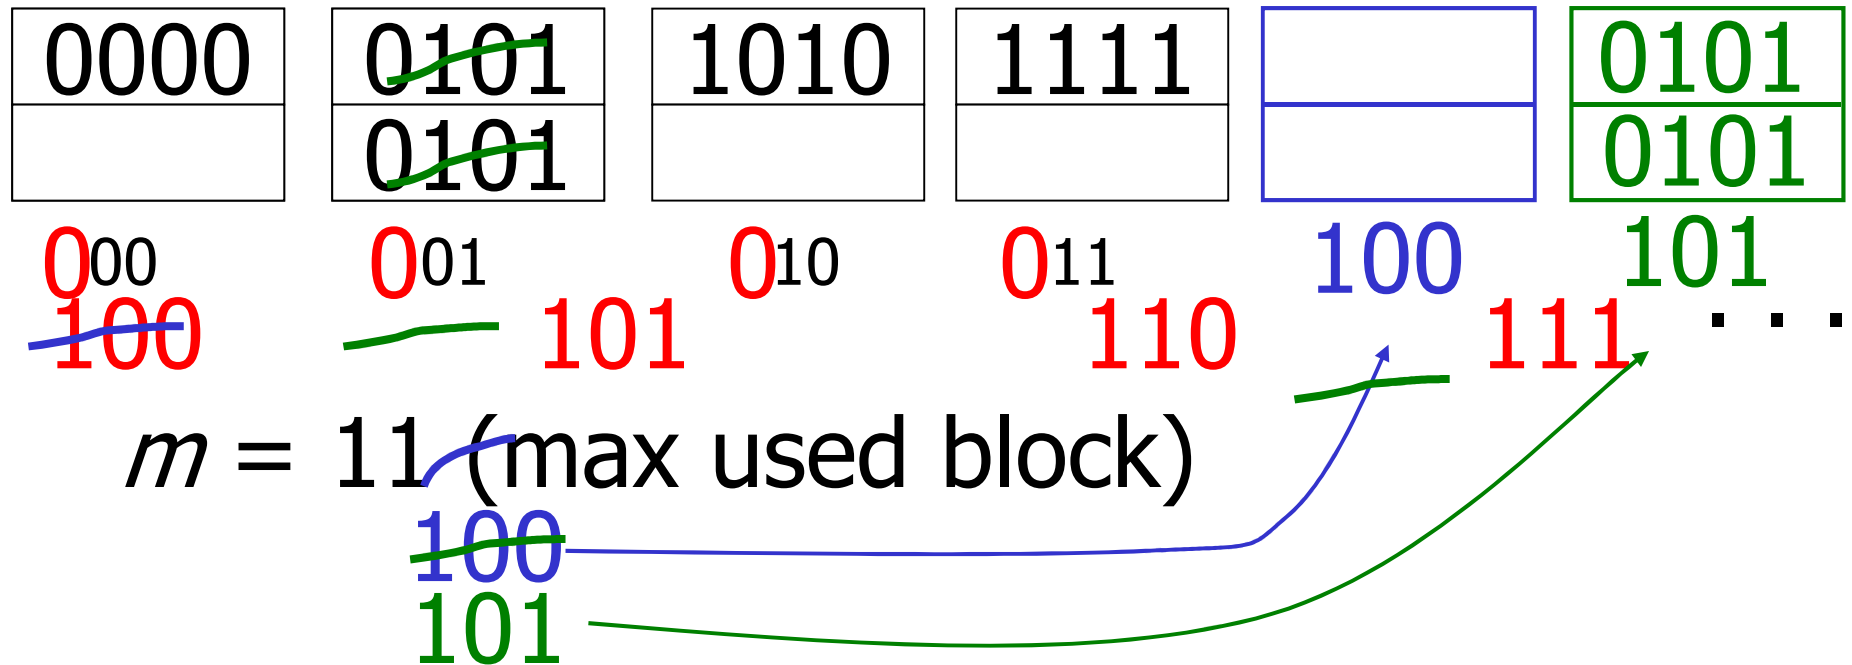
$$i = \cancel{2}3$$





Example Continued: How to grow beyond this?

$$i = \cancel{2}3$$





Summary Linear Hashing

- ⊕ Can handle growing files
 - with less wasted space
 - with no full reorganizations
- ⊕ No indirection like extensible hashing
- ⊖ Can still have overflow chains



Hash-based Indices in RDBMS:

- PostgreSQL:

`CREATE INDEX` name `ON` table `USING hash (column);`

<http://www.postgresql.org/docs/9.1/static/indexes-types.html>



Read

- Reference:

- dropbox:

<https://www.dropbox.com/s/fqv14g1zqhhl6k5/Chaps14-19.pdf>

- gdrive:

https://drive.google.com/file/d/0B03SaNyIsL_2U2pmMjVrMDI4MWs/view?usp=sharing

- Read: pg 648-661